Hyungyu Jung

Email : jhubert@snu.ac.kr

Research Interest

Accelerating LLM Inference, Parallel Computing, SW/HW Co-Design

Education

- **Seoul National University** — Seoul, South Korea
  B.S. in Electrical and Computer Engineering; GPA: 3.65/4.3 — *Mar. 2021 - Present*
  - Courses: Scalable High-Performance Computing, Embedded System Design, Digital System Design, Computer Organization, Operating System
- **Busan-Il Science High School** — Busan, South Korea
  Mar. 2019 - Feb. 2021

Projects

- **Accelerate Pipeline-Parallel LLM Inference**
  Conducted as research for my graduation thesis.
- **Accelerate Diffusion Inference on Embedded System**
  Accelerate Diffusion Inference on Nvida Jetson Orin Nano(CPU, GPU) and XILINX ZYNQ-7000(FPGA).
  Conducted as a term project for the Embedded System Design course.
- **Accelerate CNN based Text Classifier on GPU Cluster**
  Accelerate CNN based Text Classifier on GPU Cluster using CUDA.
  Conducted as a term project for the Scalable High-Performance Computing course.
- **Implementing CNN Accelerator**
  Implementing systolic-array CNN accelerator on FPGA using verilog.
  Conducted as a term project for the Digital System Design course.

Skills

- **Programming** - C, C++, Python, CUDA, Verilog
- **Frameworks** - Pytorch, vLLM

Research Experiences

- **MLSYS, Seoul National University** — Seoul, South Korea
  Undergraduate Research Intern (Advisor: Jiwon Seo) — *Mar. 2025 - Present*
- **AISys, Seoul National University** — Seoul, South Korea
  Undergraduate Research Intern (Advisor: Jinho Lee) — *Jun. 2024 - Aug. 2024*
- **AISys, Seoul National University** — Seoul, South Korea
  Undergraduate Research Intern (Advisor: Jinho Lee) — *Dec. 2023 - Feb. 2024*

Scholarship

- **Sooyoungro Church Scholarship**
  Received a merit-based full-tuition scholarship, regardless of religious affiliation. — *Mar. 2021 - Feb. 2025*